

Chapter 1: Introduction





Chapter 1: Introduction

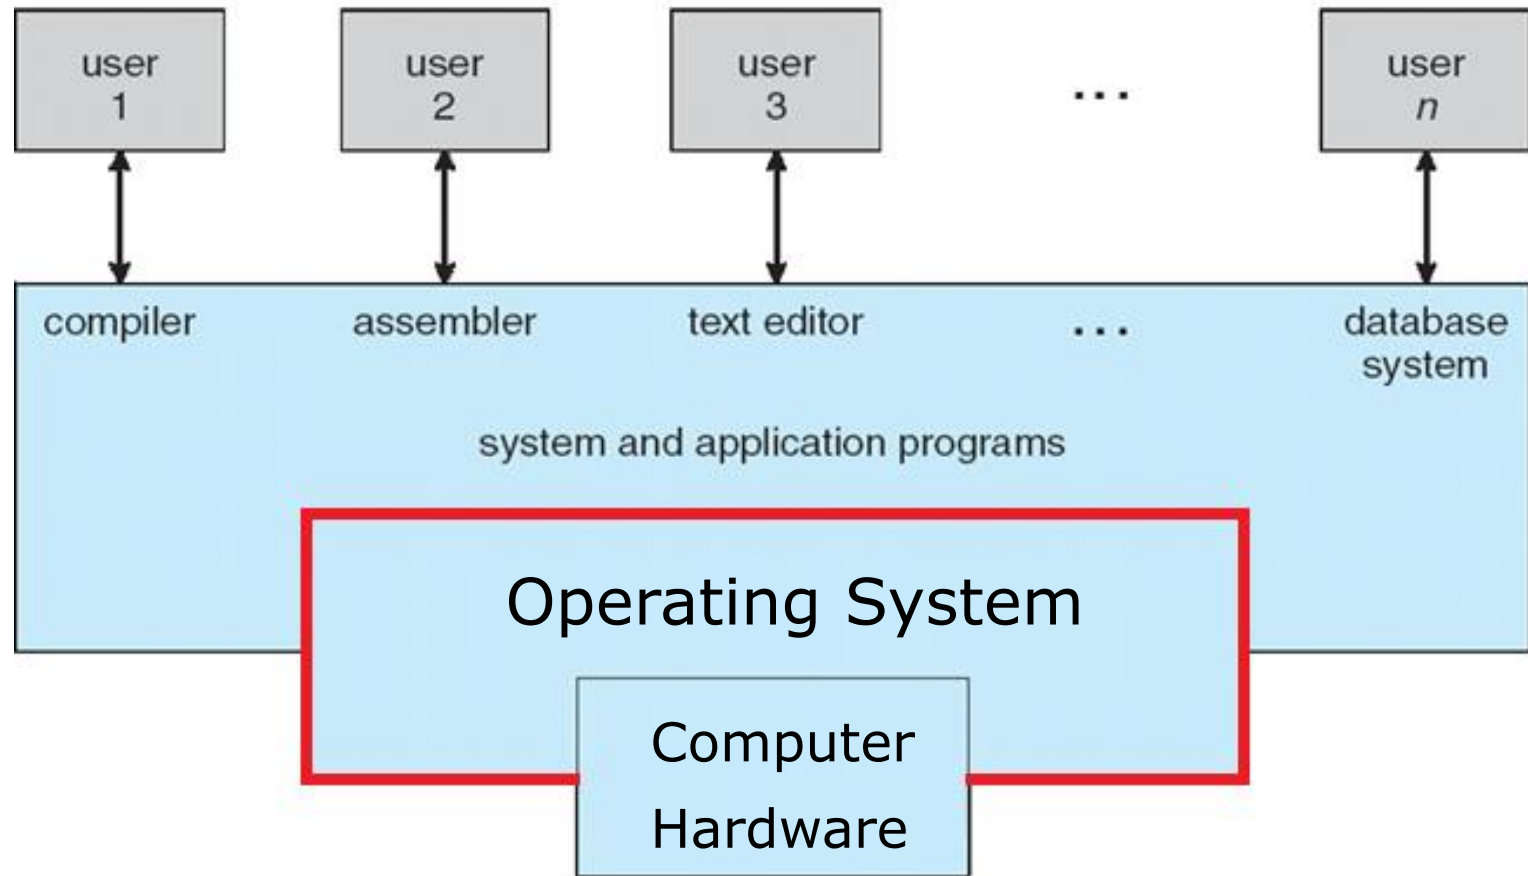
- What Operating Systems Do
- Computer-System Organization
- Operating-System Operations
 - Process Management
 - Memory Management
 - Storage Management
 - Protection and Security
- Kernel Data Structures





What Operating Systems Do

Four Components of a Computer System





What Operating Systems Do

What is an Operating System?

- A program that acts as an intermediary between a user of a computer and the computer hardware
- Operating system goals:
 - Execute user programs
 - Make the computer system convenient to use
 - Use the computer hardware in an efficient manner





What Operating Systems Do

Different Goals for Different Systems

- Depends on the point of view
- For **personal computer**, users want convenience → Care about **ease of use or performance**, but not resource utilization
- But shared computer such as **mainframe** or **minicomputer** must keep all users happy → optimize **resource utilization**
- **Handheld** computers are resource poor, optimized for **usability and battery life**
- Some computers have **little or no user interface**, such as **embedded computers** in devices and automobiles





What Operating Systems Do

Operating System Definition

- OS is a **resource allocator**
 - Manages all resources
 - Decides between conflicting requests for efficient and fair resource use
- OS is a **control program**
 - Controls execution of programs to prevent errors and improper use of the computer
- “The one program running at all times on the computer” is the **kernel**. Everything else is either a system program (ships with the operating system) or an application program.
- **No universally accepted definition**

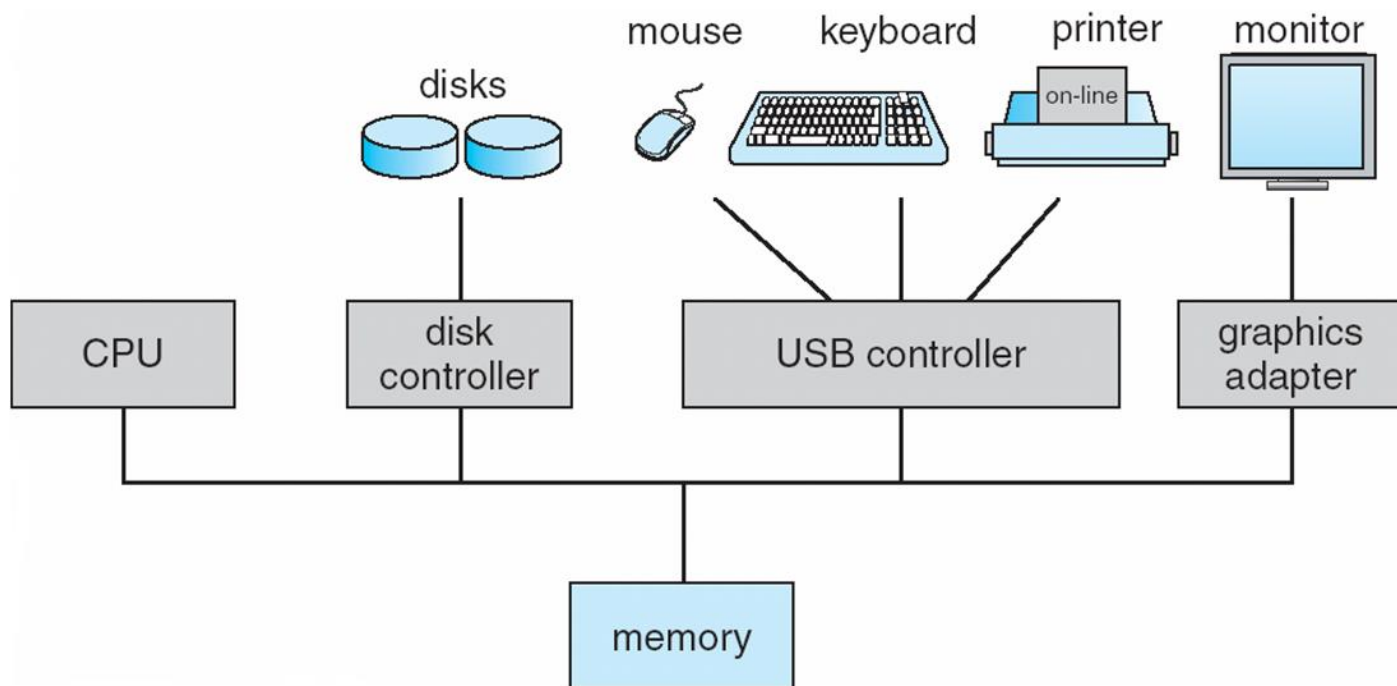




Computer System Organization

Computer-System Operation

- One or more CPUs, device controllers connect through common bus providing access to shared memory
- Concurrent execution of CPUs and devices





Computer System Organization

Computer-System Operation: Computer Startup

- **bootstrap program** is loaded at power-up or reboot
 - Typically stored in ROM or EPROM, known as **firmware**
 - Initializes all aspects of system
 - ▶ CPU registers, memory content, device controller
 - Loads operating system kernel and starts execution
- Kernel starts providing services to the system.
 - In UNIX, first system process is “init” which starts many other daemons.
- The system waits for some event to occur
 - OS is event driven (Interrupt driven)





Computer System Organization

Computer-System Operation: Interrupts

- The occurrence of an event is signaled by an **interrupt** from either hardware or software
 - Hardware can trigger an interrupt at any time by sending a **signal to CPU** (by way of system bus)
 - Software-generated interrupt is called a **trap** or **exception**, caused either by an error or a user request (executing a special operation called **system call**)
- Interrupt Service Routine (**ISR**)
 - When CPU is interrupted, it stops what it is doing and immediately transfer execution to ISR
 - On ISR completion, CPU resumes to the interrupted instruction.





Computer System Organization

Computer-System Operation: Interrupts

- Interrupt transfers control to the interrupt service routine generally, through the **interrupt vector**, which contains the addresses of all the service routines
- The operating system preserves the state of the CPU by storing **registers** and the **program counter** (the address of the interrupted instruction)
- An operating system is **interrupt driven**





Computer System Organization

Storage Structure

- Main memory – CPU can only load instructions only from memory, so any programs to run must be stored there.
 - **Random access**, typically **volatile** → e.g., SRAM, DRAM
 - non-volatile: EEPROM, ROM (bootstrap)
- Secondary storage – extension of main memory that provides large **nonvolatile** storage capacity
 - Magnetic disks – metal platters covered with magnetic material
 - ▶ Disk surface is logically divided into **tracks**, subdivided into **sectors**
 - ▶ The **disk controller** determines the logical interaction between the device and the computer
 - **Solid-state disks (SSD)** – faster than magnetic disks, nonvolatile
 - ▶ Various technologies, Becoming more popular






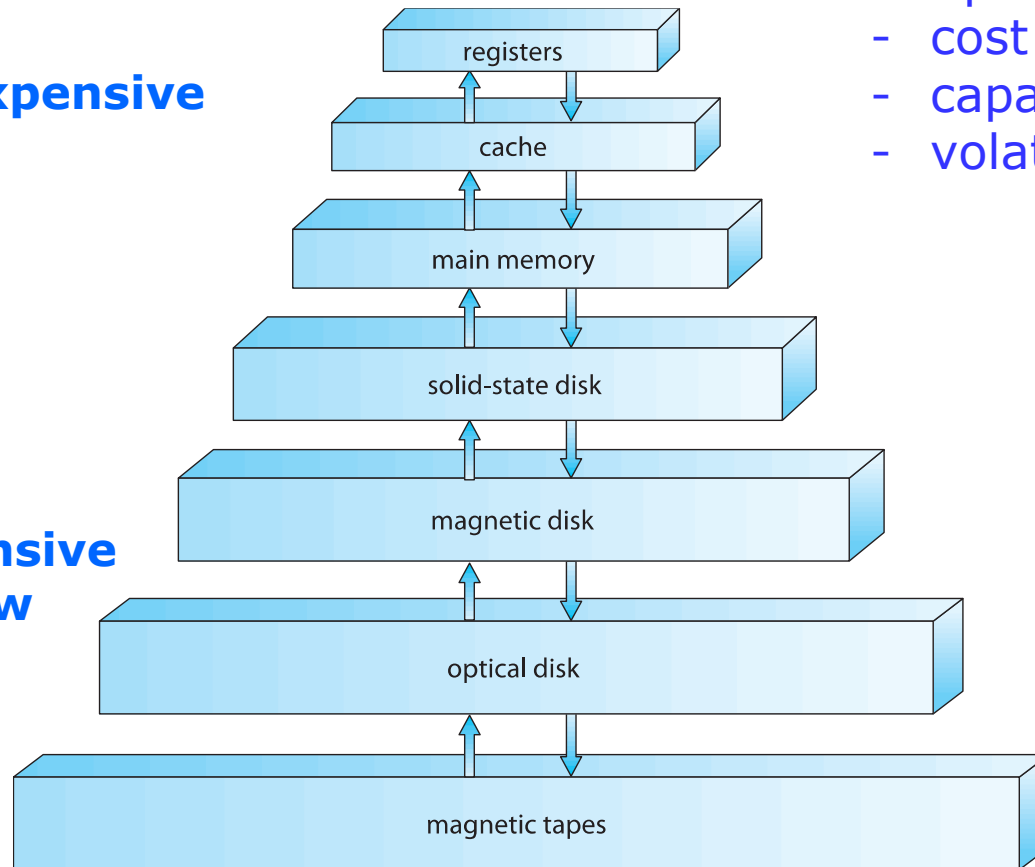
Computer System Organization

Storage Hierarchy

Fast but expensive



**Less expensive
but slow**



Difference among various storage

- Speed
- cost
- capacity
- volatility





Computer System Organization

Storage Structure - Caching

- Important principle, performed at many levels in a computer (in hardware, operating system, software)
- Information in use copied from slower to faster storage temporarily
 - e.g., main memory can be viewed as a cache for secondary storage
- Faster storage (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there
- Cache smaller than storage being cached
 - Cache management is an important design problem
 - Cache size and replacement policy






Computer System Organization

I/O Structure

- I/O devices and the CPU can execute concurrently
 - Each device controller is in charge of a particular device type
 - Each device controller has a local buffer and a set of registers
 - OS has a device driver for each device controller
 - I/O operation
 - Device driver starts an I/O operation by sending commands to the proper register of the device controller
 - Device controller examines the command register and moves data to/from the device from/to its local buffer
 - Device controller informs CPU that it has finished its operation by causing an *interrupt*
 - CPU moves data from/to main memory to/from local buffers
- ➔ Interrupt-driven I/O





Computer System Organization

I/O Structure

■ Direct Memory Access (DMA)

- Used for high-speed I/O devices able to transmit information at close to memory speeds
- Device controller transfers blocks of data from local buffer directly to main memory without CPU intervention
- Only one interrupt is generated per block, rather than the one interrupt per byte

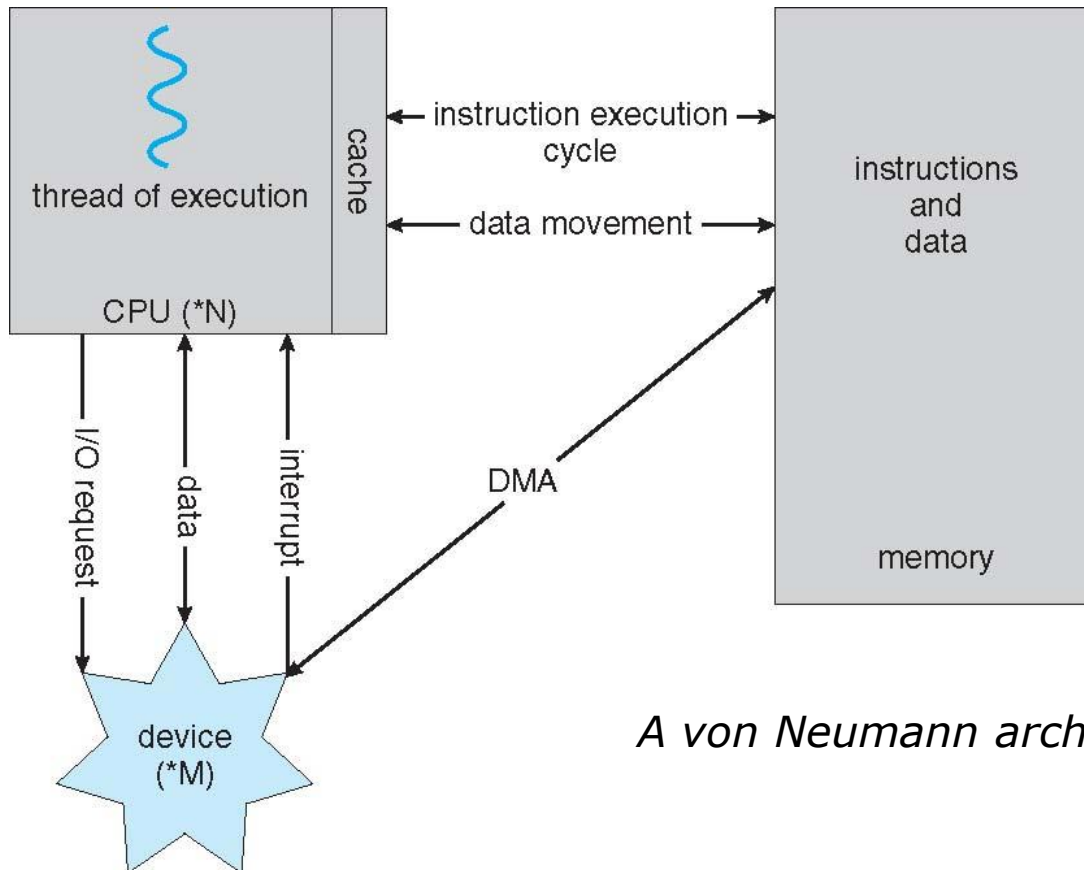




Computer-System Architecture

■ Single Processor System

- Most systems use a **single general-purpose** processor
 - ▶ Some systems have special-purpose processors as well




A von Neumann architecture





Computer-System Architecture


- **Clustered** Systems (loosely coupled)
 - Multiple systems working together by network
 - Usually sharing storage via a **storage-area network (SAN)**
 - Provides a **high-availability** service which survives failures
 - ▶ **Asymmetric clustering** has one machine in hot-standby mode
 - ▶ **Symmetric clustering** has multiple nodes running applications, monitoring each other





Computer-System Architecture

- **Multiprocessors** systems (parallel systems, tightly-coupled systems)
 - ▶ Increased throughput
 - ▶ Economy of scale
 - ▶ Increased reliability: graceful degradation (or fault tolerant)
- Two types:
 1. **Asymmetric Multiprocessing** (master-slave relationship)
 2. **Symmetric Multiprocessing**






Computer-System Architecture

- **Multiprocessors** systems : **Multicore**
 - ▶ **More efficient** than multiple chips with single cores because on-chip communication is faster than between-chip communication.
 - ▶ **Less power** than multiple single-core chips

Dual-core chip





Operating System Structure

- **Multiprogramming** (increase CPU utilization)
 - Single user cannot keep CPU and I/O devices busy at all times
 - Multiprogramming organizes jobs (code and data) so CPU always has one to execute
 - A subset of total jobs in system is kept in memory
 - One job selected and run via **job scheduling**
 - When it has to wait (for I/O for example), OS switches to another job
- **Timesharing (multitasking)** (creating interactive computing)
 - Response time should be < 1 second
 - If several jobs ready to run at the same time \Rightarrow **CPU scheduling**
 - If processes don't fit in memory, **swapping** moves them in and out
 - **Virtual memory** allows execution of processes not all in memory





Operating-System Operations

- A properly designed system must ensure that an incorrect program cannot cause other program to execute incorrectly.
- **Dual-mode** operation allows OS to protect itself and other system components
 - **User mode** and **kernel mode**
 - **Mode bit** provided by hardware
 - ▶ To distinguish when system is running user code or kernel code
 - ▶ **privileged** instructions: only used in kernel mode
 - ▶ System call changes mode to kernel, return from call resets it to user mode
- Increasingly CPUs support multi-mode operations
 - i.e. **virtual machine manager (VMM)** mode for guest **VMs**

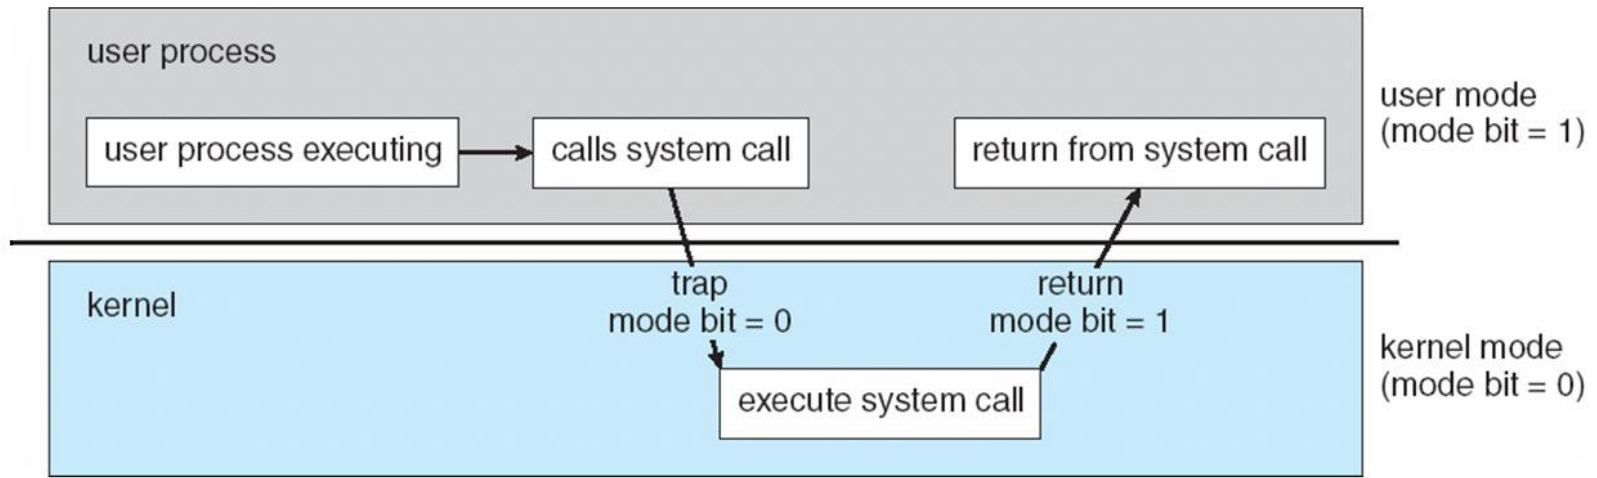




Operating-System Operations

Transition from User to Kernel Mode

- Timer to prevent infinite loop / process hogging resources
 - Set interrupt after specific period
 - Operating system decrements counter
 - When counter zero generate an interrupt
 - Set up before scheduling process to regain control





Operating-System Operations

Process Management

- A process is a program in execution. It is a unit of work within the system. Program is a ***passive entity***, process is an ***active entity***.
- Process needs resources to accomplish its task
 - CPU, memory, I/O, files
 - Initialization data
- Process termination requires reclaim of any reusable resources
- Single-threaded process has one **program counter** specifying location of next instruction to execute
 - Process executes instructions sequentially, one at a time, until completion
- Multi-threaded process has one program counter per thread
- Typically system has many processes running concurrently on one or more CPUs
 - Concurrency by multiplexing CPUs among them





Operating-System Operations

Process Management Activities

The operating system is responsible for the following activities in connection with process management:

- Creating and deleting both user and system processes
- Suspending and resuming processes
- Providing mechanisms for process synchronization
- Providing mechanisms for process communication
- Providing mechanisms for deadlock handling





Operating-System Operations

Memory Management

- All data in memory before and after processing
- All instructions in memory in order to execute
- Memory management determines what is in memory when
 - Optimizing CPU utilization and computer response to users
- Memory management activities
 - Keeping track of which parts of memory are currently being used and by whom
 - Deciding which processes (or parts thereof) and data to move into and out of memory
 - Allocating and deallocating memory space as needed





Operating-System Operations

Storage Management

- OS provides uniform, logical view of information storage
 - Various storage devices (i.e., disk drive, tape drive)
 - ▶ Varying properties include access speed, capacity, data-transfer rate, access method (sequential or random)
 - OS abstracts physical properties to logical storage unit - **file**
- File-System management
 - Files usually organized into directories
 - Access control on systems to determine who can access what
 - OS activities include
 - ▶ Creating and deleting files and directories
 - ▶ Primitives to manipulate files and dirs
 - ▶ Mapping files onto secondary storage
 - ▶ Backup files onto stable (non-volatile) storage media






Operating-System Operations

Mass-Storage Management

- Usually disks used to store data that does not fit in main memory. Entire speed of computer operation hinges on disk subsystem and its algorithms



Migration of Integer A from Disk to Register

- OS activities
 - Free-space management
 - Storage allocation
 - Disk scheduling
- Multitasking environments must be careful to use most recent value, no matter where it is stored in the storage hierarchy
- Multiprocessor environment must provide **cache coherency** in hardware such that all CPUs have most recent value in their cache





Operating-System Operations

I/O Subsystem

- One purpose of OS is to hide peculiarities of hardware devices from the user
- I/O subsystem responsible for
 - Memory management of I/O including **buffering** (storing data temporarily while it is being transferred), **caching** (storing parts of data in faster storage for performance), **spooling** (the overlapping of output of one job with input of other jobs)
 - General device-driver interface
 - Drivers for specific hardware devices





Operating-System Operations

Protection and Security


- **Protection** – any mechanism for controlling access of processes or users to resources defined by the OS
- **Security** – defense of system against internal and external attacks
 - Huge range, including denial-of-service, worms, viruses, identity theft, theft of service
- Systems generally first distinguish among users, to determine who can do what
 - User identities (**user IDs**, security IDs) include name and associated number, one per user
 - User ID then associated with all files, processes of that user to determine access control
 - Group identifier (**group ID**) allows set of users to be defined and controls managed, then also associated with each process, file
 - **Privilege escalation** allows user to change to have more rights






Kernel Data Structures


- Many similar to standard programming data structures
- ***Singly linked list***



- ***Doubly linked list***



- ***Circular linked list***





Kernel Data Structures

■ Stack

- Last in first out (LIFO)


■ Queue

- First in first out (FIFO)

■ Binary search tree

left_child \leq right_child


- **Balanced binary search tree** is $O(\lg n)$
 - ▶ *Because n items has at most $\lg n$ levels.*





Kernel Data Structures

- **Hash function** can create a **hash map**



- **Bitmap** – string of n binary digits representing the status of n items
- Linux data structures defined in ***include*** files
`<linux/list.h>`, `<linux/kfifo.h>`,
`<linux/rbtree.h>`





Open-Source Operating Systems

- Operating systems made available in source-code format rather than just binary **closed-source**
- Counter to the **copy protection** and **Digital Rights Management (DRM)** movement
- Started by **Free Software Foundation (FSF)**, which has “copyleft” **GNU Public License (GPL)**
- Examples include **GNU/Linux** and **BSD UNIX** (including core of **Mac OS X**), and many more
- Can use VMM like VMware Player (Free on Windows), Virtualbox (open source and free on many platforms - <http://www.virtualbox.com>)
 - Use to run guest operating systems for exploration

